ACCEPTED PAPER · FinNLP-65 · POSTER PRESENTATION

ADML at the FinEval 2026: Version-Aware Retrieval and Grounded Sparse Editing for Multilingual Financial QA

Agentic research across three financial NLP tasks

Gautier Marti[1] · Khaled AlNuaimi[1 2]

[1] ADIA [2] Khalifa University

1st

Greek financial NER

Entity F1 0.9676

Tied 1st

Japanese ICR

Accuracy 0.9714

2nd

Multilingual QA

ROUGE-1 F1 0.3726

Official private results · Different task metrics

Synthetic narration

AGENTIC RESEARCH METHOD

How we ran the experiments

01
Audit
Data + metric

›

02
Hypothesize
Target a failure

›

03
Implement
Controlled trial

›

04
Evaluate
Transfer + harm

›

05
Decide
Promote / reject

Evaluate across entity groups, sources, companies, and time.

Record corrections, new errors, negative results, and revised decisions.

209
research scripts

32
source modules

99
test files

25
preregistrations

Repository snapshot 557bf38 · 48,468 physical Python lines, including tests
Research code and experiment records

SYSTEMS DESIGNED FOR NEW INPUTS

Models we tried and changes we kept

Greek NER

GreekBERT · legal encoders
XLM-R · GLiNER

Numeric processing
Span consensus
Conservative repairs

Japanese ICR

BERT · Rinna RoBERTa
ModernBERT · Nemotron 9B

Domain adaptation
Ordinal decoding
Additional supervision

Multilingual QA

Eva 4B · Qwen 27B / 35B
Qwen 35B / 57B planners

Typed financial facts
Arithmetic + provenance
Grounded sparse edits

REJECTED

ModernBERT: 5/5 source folds worse

Qwen filter: one issuer exceeds error limit

A positive result: Greek consensus +4 corrections / 0 new errors (80 rows).

Greek validation repeatedly inspected; each task used its own validation protocol.

QA CASE STUDY · 192 HISTORICAL EXAMPLES

Choosing which answers to edit

ROUGE-1 F1

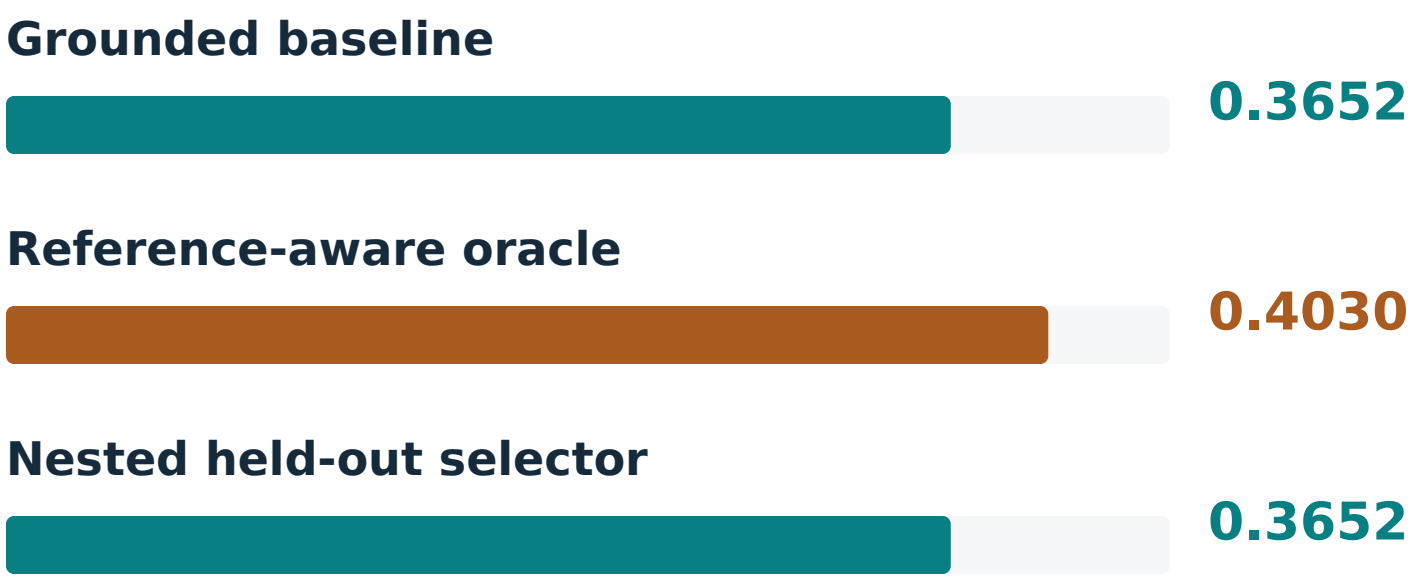

Useful edits existed.

Choosing edits for a new company remained difficult.

Nested selector: no edits.
Temporal test: 4 edits,
all 4 harmful.

Oracle: uses reference answers unavailable at inference.

Historical evaluation · Not directly comparable with the official private leaderboard.

FINAL COMPETITION SYSTEMS

What the benchmark ultimately rewarded

Greek NER

200 / 200 exact inputs

Public-reference retrieval

Japanese ICR

Explicit intent cues

Deterministic semantic rubric

Multilingual QA

76 / 76 exact inputs

Public-reference retrieval

External data permitted · Final runner: CPU only, no neural model

QA PUBLIC	QA PRIVATE
0.9994	**0.3726**
24 examples	52 examples

Same inputs. Different target versions?

Consistent with reference changes; private reference answers unavailable.

QA scores: official ROUGE-1 F1 · Retrieval coverage refers to full-query input matching.

What we learned

01 **Check data and references**

Look for repeated inputs and changed answers.

02 **Test on new sources**

Track corrections, new errors, and unchanged answers.

03 **Keep failed experiments**

Use validation results to decide what to keep.

The final benchmark left the broader model work unresolved.

Gautier Marti · Khaled AlNuaimi | **Thank you**